

# ✔ Congratulations! You passed!

Grade received 100% Latest Submission Grade 100% To pass 80% or higher

Go to next item

1. True/False: Suppose you learn a word embedding for a vocabulary of 60000 words. Then the embedding vectors could be 60000 dimensional, so as to capture the full range of variation and meaning in those words.

1 / 1 point

☒ False

☐ True

↗ Expand



**Correct**

No, the dimension of word vectors is usually smaller than the size of the vocabulary. Most common sizes for word vectors range between 50 and 1000.

2. True/False: t-SNE is a non-linear dimensionality reduction technique.

1 / 1 point

☐

False

☒

True



**Expand**



**Correct**

t-SNE is a non-linear dimensionality reduction technique.

3. Suppose you download a pre-trained word embedding which has been trained on a huge corpus of text. You then use this word embedding to train an RNN for a language task of recognizing if someone is happy from a short snippet of text, using a small training set.

<b>x (input text)</b>	<b>y (happy?)</b>
I'm feeling wonderful today!	1
I'm bummed my cat is ill.	0
Really enjoying this!	1

Then even if the word “ecstatic” does not appear in your small training set, your RNN might reasonably be expected to recognize “I’m ecstatic” as deserving a label  $y = 1$ .

☒ True

☐ False

 **Expand**

✓ **Correct**

Yes, word vectors empower your model with an incredible ability to generalize. The vector for “ecstatic” would contain a positive/happy connotation which will probably make your model classify the sentence as a "1".

4. Which of these equations do you think should hold for a good word embedding? (Check all that apply)

☒  $e_{\text{boy}} - e_{\text{brother}} \approx e_{\text{girl}} - e_{\text{sister}}$



Correct

Yes!

☐  $e_{\text{boy}} - e_{\text{girl}} \approx e_{\text{sister}} - e_{\text{brother}}$

☒  $e_{\text{boy}} - e_{\text{girl}} \approx e_{\text{brother}} - e_{\text{sister}}$



Correct

Yes!

☐  $e_{\text{boy}} - e_{\text{brother}} \approx e_{\text{sister}} - e_{\text{girl}}$

 **Expand**



**Correct**

Great, you got all the right answers.

5. Let  $E$  be an embedding matrix, and let  $o_{1234}$  be a one-hot vector corresponding to word 1234. Then to get the embedding of word 1234, why don't we call  $E * o_{1234}$  in Python?

**1 / 1 point**

- ☐ None of the above: calling the Python snippet as described above is fine.
- ☐ This doesn't handle unknown words (<UNK>).
- ☒ It is computationally wasteful.
- ☐ The correct formula is  $E^T * o_{1234}$

 **Expand**



**Correct**

Yes, the element-wise multiplication will be extremely inefficient.

6. When learning word embeddings, we create an artificial task of estimating  $P(\text{target} \mid \text{context})$ . It is okay if we do poorly on this artificial prediction task; the more important by-product of this task is that we learn a useful set of word embeddings.

**1 / 1 point**

☐ False

☒ True

 Expand

 Correct

7. In the word2vec algorithm, you estimate  $P(t \mid c)$ , where  $t$  is the target word and  $c$  is a context word. How are  $t$  and  $c$  chosen from the training set? Pick the best answer.

1 / 1 point

- ☐  $c$  is a sequence of several words immediately before  $t$
- ☐  $c$  is the sequence of all the words in the sentence before  $t$
- ☒  $c$  and  $t$  are chosen to be nearby words.
- ☐  $c$  is the one word that comes immediately before  $t$



 Expand

 Correct

8. Suppose you have a 10000 word vocabulary, and are learning 100-dimensional word embeddings. The word2vec model uses the following softmax function:

$$P(t \mid c) = \frac{e^{\theta_t^T e_c}}{\sum_{t'=1}^{10000} e^{\theta_{t'}^T e_c}}$$

True/False: After training, we should expect  $\theta_t$  to be very close to  $e_c$  when  $t$  and  $c$  are the same word.

☒ False

☐ True

1 / 1 point

 **Expand**



**Correct**

To review this concept watch the lecture.

9. Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings. The GloVe model minimizes this objective:

**1 / 1 point**

$$\min \sum_{i=1}^{10,000} \sum_{j=1}^{10,000} f(X_{ij})(\theta_i^T e_j + b_i + b_j' - \log X_{ij})^2$$

True/False:  $X_{ij}$  is the number of times word j appears in the context of word i.

☒ True

☐ False

 **Expand**



**Correct**

$X_{ij}$  is the number of times word j appears in the context of word i.

10. You have trained word embeddings using a text dataset of  $t_1$  words. You are considering using these word embeddings for a language task, for which you have a separate labeled dataset of  $t_2$  words. Keeping in mind that using word embeddings is a form of transfer learning, under which of these circumstances would you expect the word embeddings to be helpful?

- ☐ When  $t_1$  is equal to  $t_2$
- ☐ When  $t_1$  is smaller than  $t_2$
- ☒ When  $t_1$  is larger than  $t_2$

[↗ Expand](#)**Correct**

Transfer embeddings to new tasks with smaller training sets.